

Handwritten digits classification using neural network

In this notebook we will classify handwritten digits using a simple neural network which has only input and output layers. We will then add a hidden layer and see how the performance of the model improves

```
In [2]: import tensorflow as tf  
from tensorflow import keras  
import matplotlib.pyplot as plt  
%matplotlib inline  
import numpy as np
```

```
In [3]: (X_train, y_train), (X_test, y_test) = keras.datasets.mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz (https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz)  
11490434/11490434 [=====] - 1s 0us/step
```

```
In [4]: len(X_train)
```

```
Out[4]: 60000
```

```
In [5]: len(X_test)
```

```
Out[5]: 10000
```

```
In [6]: X_train[0].shape
```

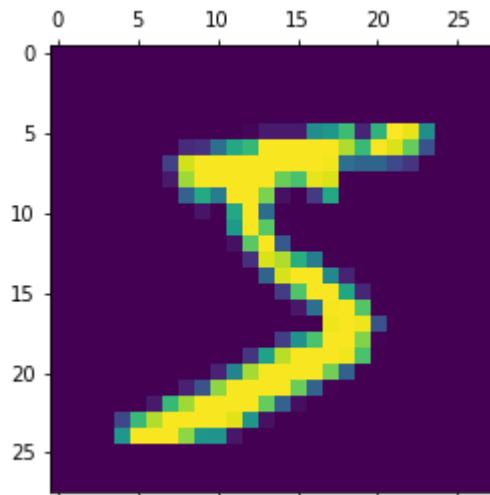
```
Out[6]: (28, 28)
```

```
In [7]: x_train[0]
```

```
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0, 46, 130, 183, 253, 253, 207,  2,  0,  0,  0,  0,  0,  0,
  0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 39,
 148, 229, 253, 253, 253, 250, 182,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 24, 114, 221,
 253, 253, 253, 253, 201,  78,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 66, 213, 253, 253,
 253, 253, 198,  81,  2,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 23, 66, 213, 253, 253,
 195, 80,  9,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 253, 253, 253, 253,
 11,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 253, 253, 253, 253,
 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 212, 135, 132, 16,
 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 0,  0,  0,  0,
 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 0,  0,  0,  0,
 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 0,  0,  0,  0,
 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 0,  0,  0,  0,
 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
 0], dtype=uint8)
```

In [8]: `plt.matshow(X_train[0])`

Out[8]: <matplotlib.image.AxesImage at 0x20591f11700>



In [9]: `y_train[0]`

Out[9]: 5

```
In [10]: X_train = X_train / 255  
X_test = X_test / 255
```

In [11]: x_train[0]

```
In [12]: X_train_flattened = X_train.reshape(len(X_train), 28*28)
        X_test_flattened = X_test.reshape(len(X_test), 28*28)
```

```
In [13]: X_train_flattened.shape
```

Out[13]: (60000, 784)

```
In [14]: x_train_flattened[0]
```

0.00392157	0.60392157	0.99215686	0.35294118	0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.99215686	0.74509804	0.00784314	0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.2745098	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.	, 0.	, 0.	, 0.
0.	, 0.1372549	, 0.94509804	0.88235294	0.62745098,
0.42352941	0.00392157	0.	, 0.	, 0.
0.	0.	0.	0.	0.

Very simple neural network with no hidden layers



```
In [15]: model = keras.Sequential([
    keras.layers.Dense(10, input_shape=(784,), activation='sigmoid')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(X_train_flattened, y_train, epochs=5)
```

```
Epoch 1/5
1875/1875 [=====] - 2s 680us/step - loss: 0.4709 - accuracy: 0.8758
Epoch 2/5
1875/1875 [=====] - 1s 652us/step - loss: 0.3037 - accuracy: 0.9149
Epoch 3/5
1875/1875 [=====] - 1s 649us/step - loss: 0.2827 - accuracy: 0.9216
Epoch 4/5
1875/1875 [=====] - 1s 647us/step - loss: 0.2730 - accuracy: 0.9240
Epoch 5/5
1875/1875 [=====] - 1s 637us/step - loss: 0.2670 - accuracy: 0.9255
```

Out[15]: <keras.callbacks.History at 0x2059194ce80>

In [16]: `model.evaluate(X_test_flattened, y_test)`

```
313/313 [=====] - 0s 563us/step - loss: 0.2715 - accuracy: 0.9245
```

Out[16]: [0.2714834213256836, 0.9244999885559082]

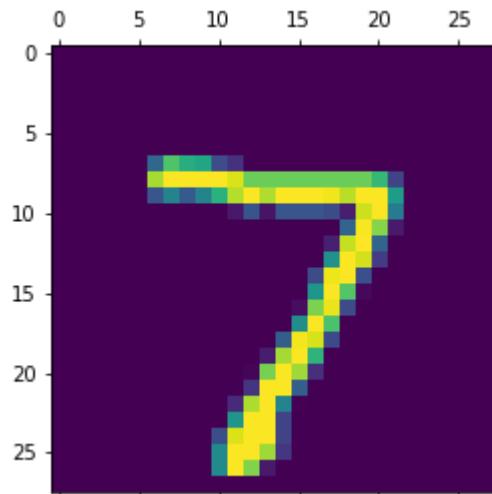
In [17]: `y_predicted = model.predict(X_test_flattened)`
`y_predicted[0]`

```
313/313 [=====] - 0s 470us/step
```

Out[17]: `array([3.68300751e-02, 3.70372277e-07, 9.22595412e-02, 9.49832320e-01,`
`1.59104867e-03, 1.35646909e-01, 1.48944628e-06, 9.99841928e-01,`
`1.00577824e-01, 6.78470373e-01], dtype=float32)`

In [18]: `plt.matshow(X_test[0])`

Out[18]: <matplotlib.image.AxesImage at 0x205c3a24ac0>



np.argmax finds a maximum element from an array and returns the index of it

In [19]: `np.argmax(y_predicted[0])`

Out[19]: 7

In [20]: `y_predicted_labels = [np.argmax(i) for i in y_predicted]`

In [21]: `y_predicted_labels[:5]`

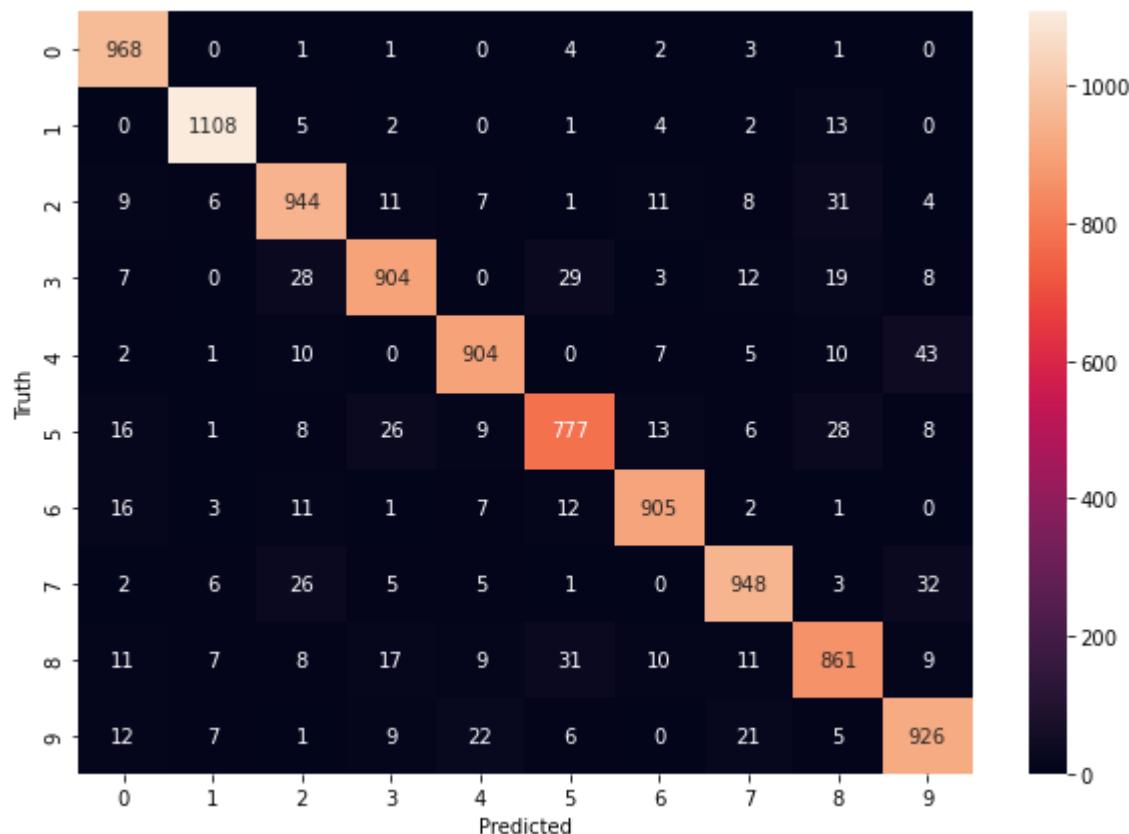
Out[21]: [7, 2, 1, 0, 4]

```
In [22]: cm = tf.math.confusion_matrix(labels=y_test,predictions=y_predicted_labels)
cm
```

```
Out[22]: <tf.Tensor: shape=(10, 10), dtype=int32, numpy=
array([[ 968,      0,      1,      1,      0,      4,      2,      3,      1,      0],
       [      0, 1108,      5,      2,      0,      1,      4,      2,     13,      0],
       [      9,      6, 944,     11,      7,      1,     11,      8,     31,      4],
       [      7,      0,     28, 904,      0,     29,      3,     12,     19,      8],
       [      2,      1,     10,      0, 904,      0,      7,      5,     10,     43],
       [     16,      1,      8,     26,      9,    777,     13,      6,     28,      8],
       [     16,      3,     11,      1,      7,     12,   905,      2,      1,      0],
       [      2,      6,     26,      5,      5,      1,      0,   948,      3,     32],
       [     11,      7,      8,     17,      9,     31,     10,     11,   861,      9],
       [     12,      7,      1,      9,    22,      6,      0,     21,      5,    926]])>
```

```
In [23]: import seaborn as sn
plt.figure(figsize = (10,7))
sn.heatmap(cm, annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

```
Out[23]: Text(69.0, 0.5, 'Truth')
```



Using hidden layer

```
In [24]: model = keras.Sequential([
    keras.layers.Dense(100, input_shape=(784,), activation='relu'),
    keras.layers.Dense(10, activation='sigmoid')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(X_train_flattened, y_train, epochs=5)
```

```
Epoch 1/5
1875/1875 [=====] - 2s 815us/step - loss: 0.2671 - accuracy: 0.9237
Epoch 2/5
1875/1875 [=====] - 2s 810us/step - loss: 0.1229 - accuracy: 0.9637
Epoch 3/5
1875/1875 [=====] - 2s 827us/step - loss: 0.0850 - accuracy: 0.9745
Epoch 4/5
1875/1875 [=====] - 2s 805us/step - loss: 0.0642 - accuracy: 0.9810
Epoch 5/5
1875/1875 [=====] - 2s 809us/step - loss: 0.0517 - accuracy: 0.9840
```

```
Out[24]: <keras.callbacks.History at 0x205c3b51ee0>
```

```
In [25]: model.evaluate(X_test_flattened,y_test)
```

```
313/313 [=====] - 0s 681us/step - loss: 0.0745 - accuracy: 0.9762
```

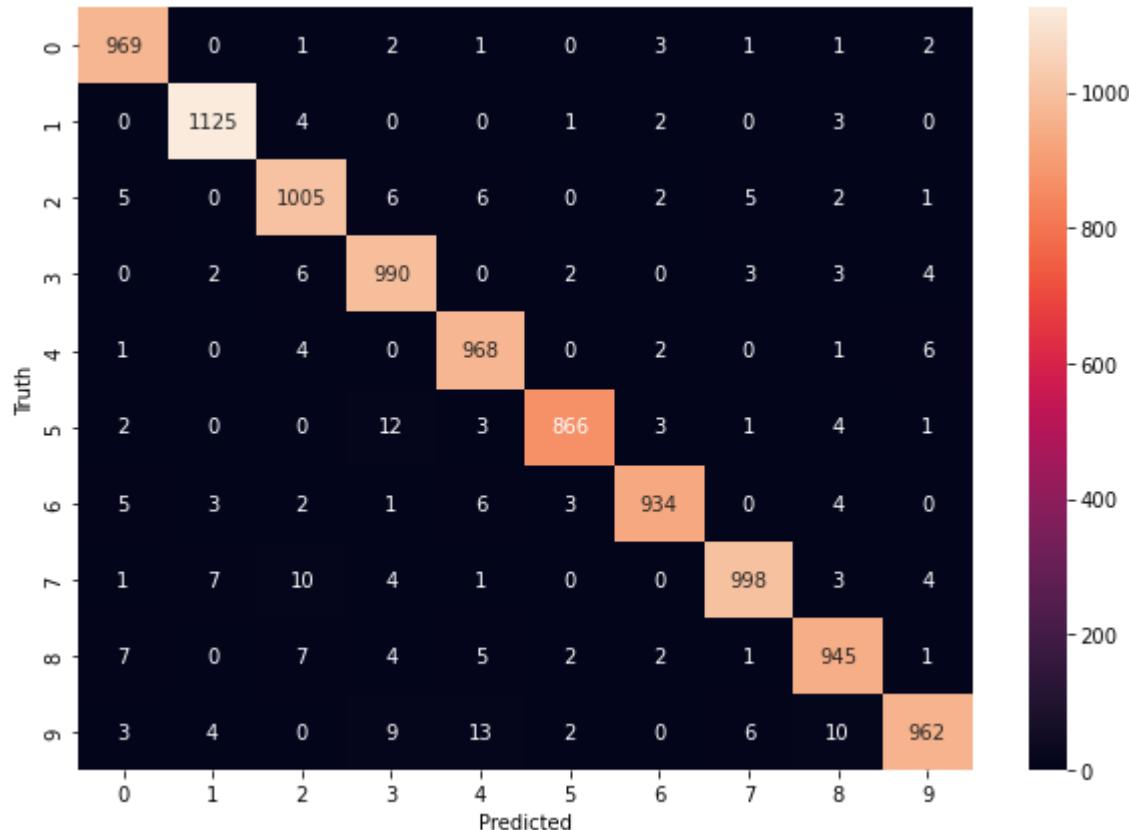
```
Out[25]: [0.07447880506515503, 0.9761999845504761]
```

```
In [26]: y_predicted = model.predict(X_test_flattened)
y_predicted_labels = [np.argmax(i) for i in y_predicted]
cm = tf.math.confusion_matrix(labels=y_test,predictions=y_predicted_labels)

plt.figure(figsize = (10,7))
sn.heatmap(cm, annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

313/313 [=====] - 0s 591us/step

Out[26]: Text(69.0, 0.5, 'Truth')



Using Flatten layer so that we don't have to call .reshape on input dataset

```
In [27]: model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(100, activation='relu'),
    keras.layers.Dense(10, activation='sigmoid')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(X_train, y_train, epochs=10)

100/100 [=====] - 2s 814us/step - loss: 0.1258 - accuracy: 0.9632
Epoch 3/10
1875/1875 [=====] - 2s 812us/step - loss: 0.0870 - accuracy: 0.9743
Epoch 4/10
1875/1875 [=====] - 2s 845us/step - loss: 0.0664 - accuracy: 0.9798
Epoch 5/10
1875/1875 [=====] - 2s 811us/step - loss: 0.0528 - accuracy: 0.9836
Epoch 6/10
1875/1875 [=====] - 2s 814us/step - loss: 0.0418 - accuracy: 0.9872
Epoch 7/10
1875/1875 [=====] - 2s 818us/step - loss: 0.0338 - accuracy: 0.9895
Epoch 8/10
1875/1875 [=====] - 2s 810us/step - loss: 0.0276 - accuracy: 0.9914
```

```
In [28]: model.evaluate(X_test,y_test)
```

```
313/313 [=====] - 0s 655us/step - loss: 0.0868 - accuracy: 0.9766
```

```
Out[28]: [0.0867568701505661, 0.9765999913215637]
```